



# Processes and Threads

# Components/Context of a Process


- Address Space
  - Code, Data, Stack, Heap
- Registers + SP + PC
- OS information: Open files, accounting info, ...
- Page table of this process tracks address space
- OS keeps other information in a Process Control Block (PCB)

# Address Space



# Context Switch

To improve efficiency: Keep contexts of several processes in memory and switch periodically between them.



```
Context_switch() {
```

```
    push R0, R1, ...           // save regs on its stack
    PCB[curr].SP = SP          // save stack pointer
    PCB[curr].PT = PT          // save ptr(s) to address space
    // NOTE: no point saving PC!
```

```
    next = schedule()          // find next process to run
```

```
    PT = PCB[next].PT
    SP = PCB[next].SP
    pop Rn, ... R0
```

```
    return                     // NOTE: Ctrl returns to another process
}
```

# Processes are convenient to ...

- Implement concurrency (between users, between activities of a user, ...)
- Insulate one activity from another

# Processes are NOT desirable because ...

- They are heavy weight - higher scheduling (context switch) costs
  - Direct costs of switching address spaces.
  - Indirect costs (e.g. TLB/cache flushes)
- State Sharing is a problem
  - System V extensions to address this problem.

# Solution: Threads

- Activities within the same address space.
- Threads within a process share (code, data, heap).
- Only stacks are disjoint.
- Switching between threads only involves switching stacks.
- Sharing is implicit
- No protection between threads of a process - but this is OK since they are meant to be cooperative.



# A Disk Server: Using 1 thread of control

```

Server() {
  while () {
    receive request;
    case (REQUEST) {
      READ:
        If device busy {
          queue request
        }
        else {
          send request to disk
          queue for response
        }
        .....
      RESPONSE_FROM_DISK:
        check concerned queue
        Dequeue
        Reply back to requester
    }
  }
}

```

Cannot block!  
Else you will  
lose  
concurrency!

# A Disk Server: Using Multiple Threads

```
Server() {
    while () {
        receive_request()
        case (REQUEST) {
            READ: fork_thread(read_disk, ....);
            WRITE: fork_thread(write_disk,...);
            ...
        }
    }
}
```

```
Read_disk() {
    While (device busy)
        ;
    Program disk
    While (results not available)
        ;
    Reply back to requester
}
```

Can afford to  
block without loss  
in throughput

# Advantages of threads

- Cheaper to switch than processes
- Easier to program (do not worry about blocking)
- Easier to communicate (memory is implicitly shared)

# Threads can be implemented on

- (a) Uniprocessors
- (b) Symmetric Multiprocessors (SMPs)
- (c) Cache coherent NUMA Multiprocessors (CC-NUMA)
- (d) Message passing multiprocessors
- (e) Networked/distributed workstations

- Efficiency decreases from (b) to (e)

# Flavors of threads

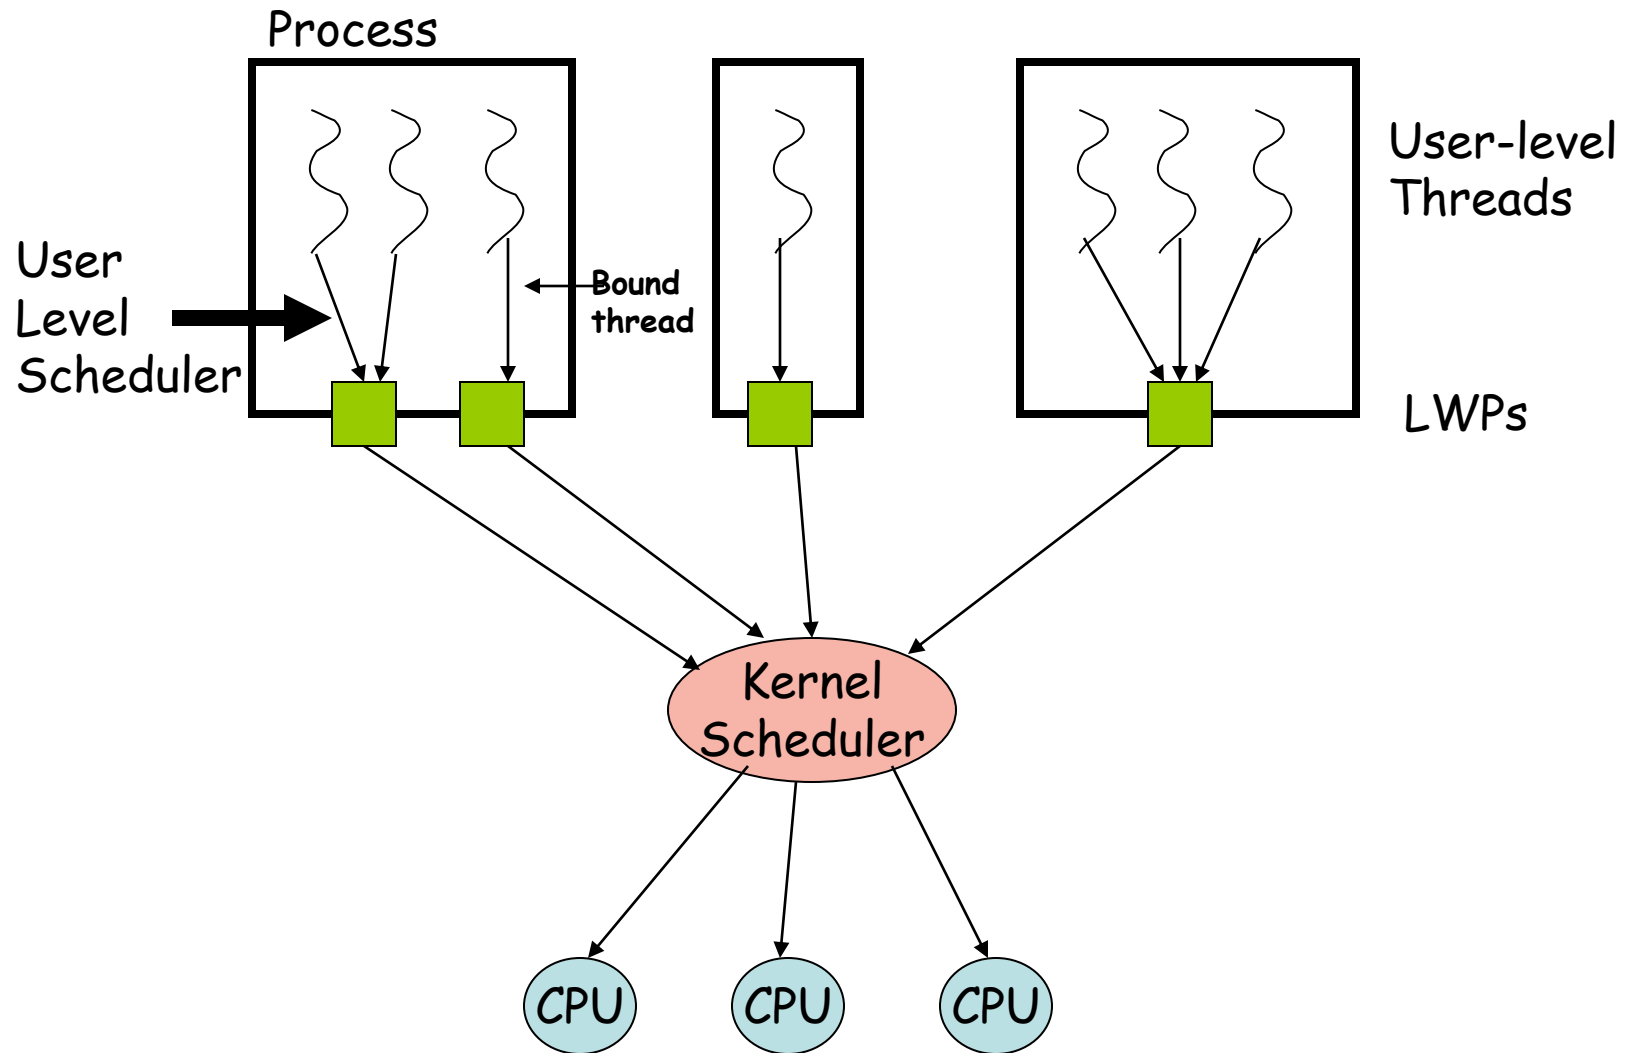
- Entirely user-level threads
  - Threads are created and managed entirely by the user code/libraries
  - Kernel is unaware of their existence
  - Advantages: Scheduling/switching can be more efficient
- Kernel-level threads
  - Kernel is the one creating (by explicit user calls) and managing the threads
  - Advantages: Better resource allocation across address spaces.

## Solaris Multithreaded Architecture

- Both user and kernel level threads
- Kernel level threads are synonymous with Light Weight Processes (LWPs)
- In addition, user can create user-level threads that will be run by these LWPs


- User is provided with processes
- User can create user-level threads via a user-level library (e.g. pthreads)
- User can specify how many LWPs should run these user-level threads
- User can bind threads to LWPs if needed or can dynamically schedule the threads to LWPs (via the library).
- Kernel schedules the LWPs on the available CPUs.

# Solaris Architecture

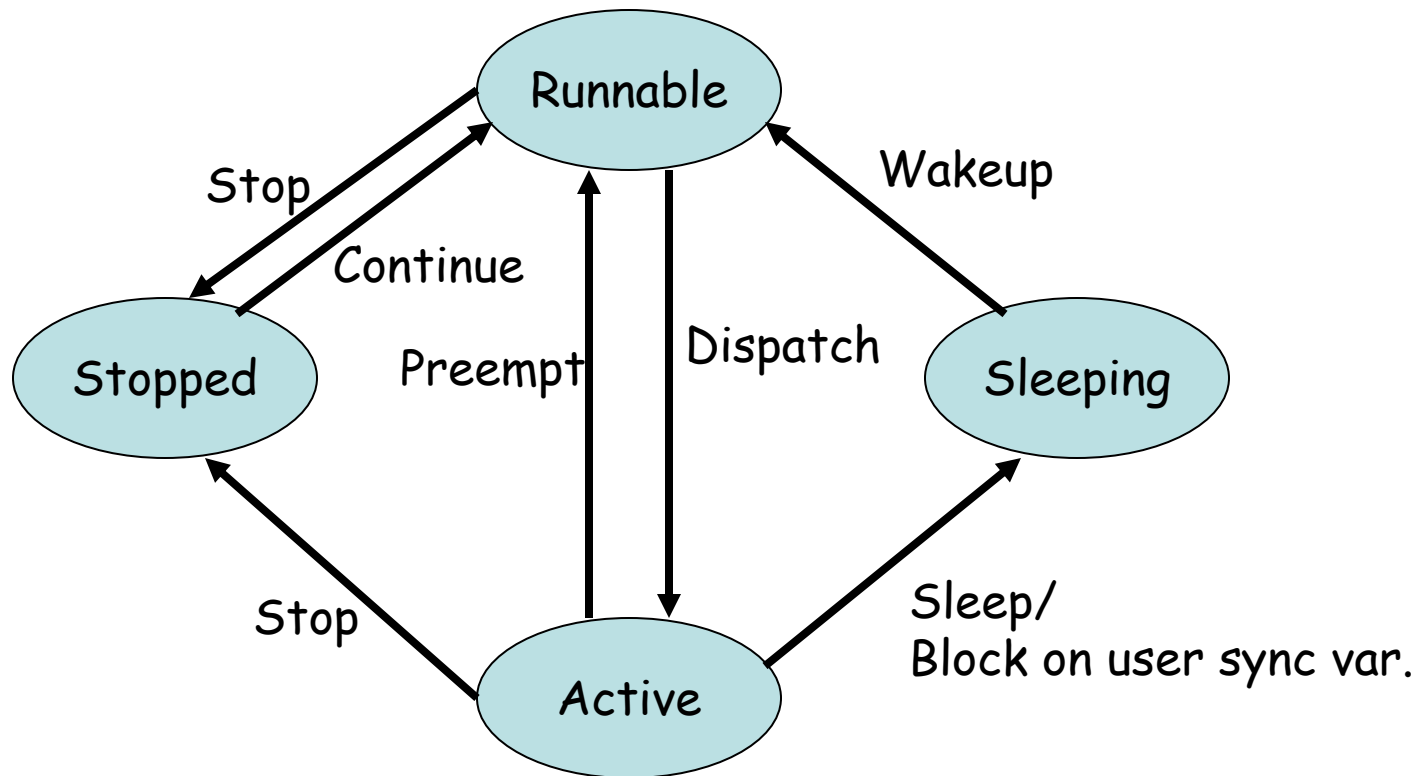




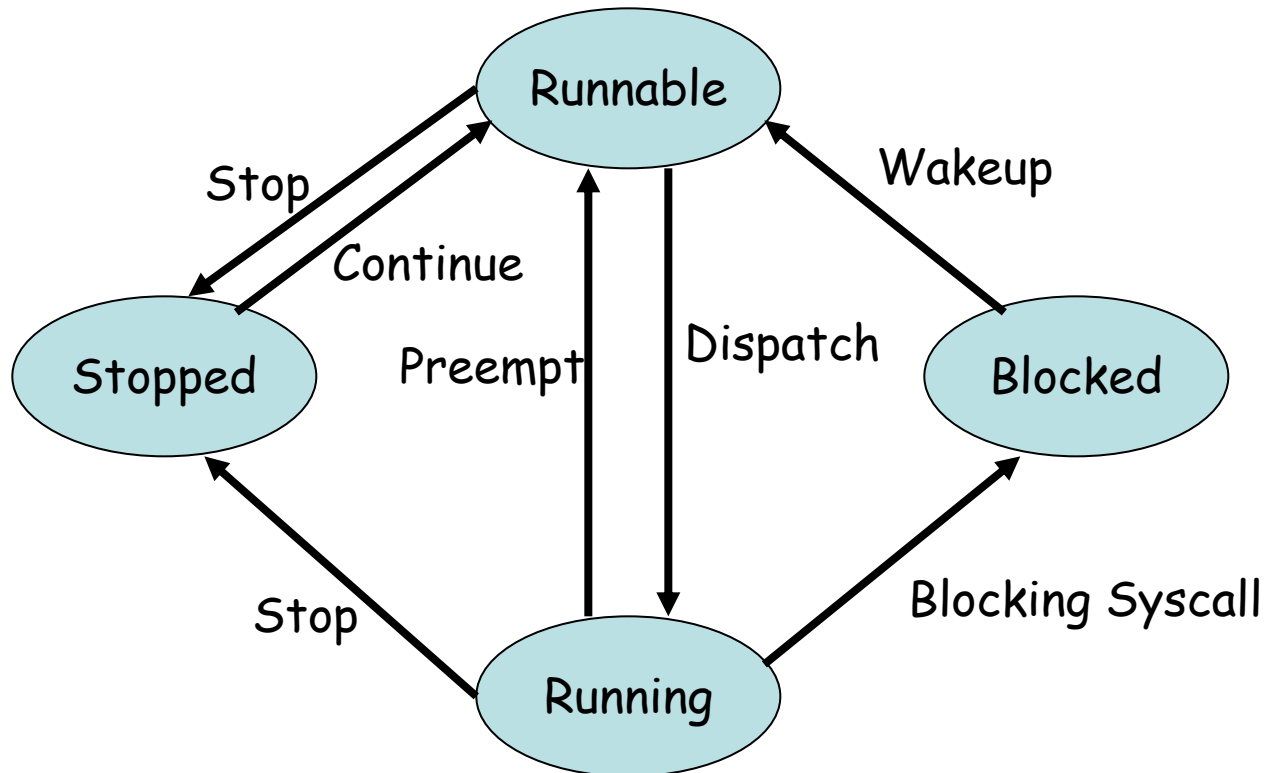

# User-level scheduling



# User-level thread States



# LWP States



# SIGWAITING

- When all LWPs of a process are blocked in the kernel, the process is sent a SIGWAITING signal.
- You can choose to add create LWPs if needed
- This is a way of avoiding the problem of the kernel being ignorant of the workload within the process.

# User-level scheduling

- Can be pre-emptive or not
- You can implement your own scheduler

# Scheduling LWPs

- 3 scheduling classes
  - Timesharing (TS)
    - For user processes
    - Time slices the CPU
    - 0-59 priority levels
  - System
    - For some kernel activity (interrupts, etc.)
    - Priority levels 60-99
    - Usually fixed priority
  - Real-time
    - Very time sensitive
    - Fixed (higher) priorities 100-159

# Threads Interface

- Thread creation, termination
- Communication (implicitly shared memory)
- Synchronization (mutex, cond variables, semaphores, etc.)

# Thread Control

- `t = t_create(func, arg, flags)`: Creates a thread that would start executing `func()`
- `t_exit()`: thread is done
- `t_join(t)` waits for the specified thread to exit
- `t_set_prio()/t_get_prio()`: set/get thread priorities
- `t_yield()`: voluntarily relinquishing CPU



# Synchronization

- Locks
  - Condition Variables
  - Semaphores
  - Monitors
- 
- What determines your choice of which sync mechanism to use?

# Locks

- Used to guard sections of code where shared data is manipulated
- Ordering is not as important (only exclusion)
- Waiting for events to happen is not as easy to implement.
- Exercise: How would you implement a lock construct (lock/unlock operations) in a threads library?

# Condition Variables

- `C_wait()` and `c_signal()` operations.
- A thread blocked on `c_wait()` returns when another performs a `c_signal()`.
- What differentiates a condition variable from a (boolean) semaphore?
- Signals can get lost!
  - i.e. if the signal is done before the wait, then signal is lost

```
Cond_t not_full, not_empty;;
Int count == 0;
```

```
Append() {
    if count == N    c_wait(not_full);

    ... ADD TO BUFFER, UPDATE COUNT ...

    c_signal(not_empty);
}
```

```
Remove() {
    if count == 0    c_wait(not_empty);

    ... REMOVE FROM BUFFER, UPDATE COUNT

    c_signal(not_full);
}
```

What is wrong?

```

Cond_t not_full, not_empty;
Mutex_lock m;
Int count == 0;

Append() {
    mutex_lock(m);
    if count == N    c_wait(not_full,m);

    ... ADD TO BUFFER, UPDATE COUNT ...

    c_signal(not_empty);
    mutex_unlock(m);
}

Remove() {
    mutex_lock(m);
    if count == 0    c_wait(not_empty,m);

    ... REMOVE FROM BUFFER, UPDATE COUNT

    c_signal(not_full);
    mutex_unlock(m);
}

```

NOTE: You can improve this code for more concurrency!

# Semaphores

- Condition variables with state (signals) preserved
- Counting semaphores offer more flexibility
- P() and V() operations.

# Monitors

- Encapsulation of Shared Data Objects and operations on them (an ADT), with the property that only 1 activity can be “virtually running” (i.e. running+ready) at any time.
- Operations use synchronization so that the invoker does not need to worry about them.
- You can have wait/signal primitives within the monitor to block/proceed when needed (note they are similar to condition variables, except you do not need a lock to go with it!).

```
Monitor Bounder_Buffer {
```

```
    Buffer[0..N-1];
```

```
    Int count= 0, head=tail=0;
```

```
    Cond_t not_full, not_empty;
```

```
    Append(Data) {
```

```
        if count == N wait(not_full);
```

```
        Buffer[head] = Data
```

```
        count++;
```

```
        head = (head+1)%N;
```

```
        if !empty(not_empty) signal(not_empty);
```

```
    }
```

```
    Remove() {
```

```
        if count == 0 wait(not_empty);
```

```
        Data = Buffer[tail];
```

```
        count--;
```

```
        tail = (tail+1)%N;
```

```
        if !empty(not_full) signal(not_full);
```

```
    }
```

```
}
```



# Monitor Readers\_Writers {

```
Database [];
Int nreaders, nwriters;
Cond readq, writeq;
```

```
Read() {
    while (nwriters > 0) wait(readq);
    nreaders++;
    signal(readq);

    .... Read database ...

    nreaders--;
    if (nreaders == 0) signal(writeq);
}

Write() {
    while (nreaders > 0) || (nwriters > 0) wait(writeq);
    nwriters++;

    ... Write database

    nwriters--;
    if !empty(readq) signal(readq);
    else if !empty(writeq) signal(writeq)
}
}
```

# Monitor Readers\_Writers {

```
Int nreaders, nwriters;
```

```
Cond readq, writeq;
```

```
Start_Read() {
```

```
    while (nwriters > 0) wait(readq);
```

```
    nreaders++;
```

```
    signal(readq);
```

```
}
```

```
End_Read() {
```

```
    nreaders--;
```

```
    if (nreaders == 0) signal(writeq);
```

```
}
```

```
Start_Write() {
```

```
    while (nreaders>0) || (nwriters>0) wait(writeq);
```

```
    nwriters++;
```

```
}
```

```
End_Write() {
```

```
    nwriters--;
```

```
    if !empty(readq) signal(readq);
```

```
    else if !empty(writeq) signal(writeq)
```

```
}
```

```
}
```